

Homework 1

Kendall Brown 8564403

Fall 2018

```
algae = read.table("algae.txt", header=T, na.strings="NA")
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##   filter, lag
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

labc. Counting observations of algae and comparing mean v median and variance v MAD.

#1a.

```
Datasum=summarise(algae,count=n())
Datasum##number of obs in algae data set
```

```
##   count
## 1    180
```

#1b.

```
?summarise
?summarise_all
?summarise_at
w=algae%>%summarise_at(c(6:11),mean,na.rm=T)
w
```

```
##      Cl      NO3      NH4      oP04      P04      Chla
## 1 40.2932 2.946356 414.1545 75.24621 133.2588 10.62003
```

```
x=algae%>%summarise_at(c(6:11),var,na.rm=T)
x
```

```
##      Cl      NO3      NH4      oP04      P04      Chla
## 1 1797.325 5.40628 1240777 9421.302 18136.13 238.7067
```

```
y=algae%>%summarise_at(c(6:11),median,na.rm=T)
y
```

```
##      Cl      NO3      NH4      oP04      P04      Chla
## 1 29.5 2.262 103.0415 35.928 85.95 4.5
```

```
z=algae%>%summarise_at(c(6:11),mad,na.rm=T)
z
```

```
##      Cl      NO3      NH4      oP04      P04      Chla
## 1 31.8759 2.095655 115.0868 44.37125 102.9206 5.355892
```

```
w>y ##Test if mean > median
```

```
##      Cl      NO3      NH4      oP04      P04      Chla
```

```
## [1,] TRUE TRUE TRUE TRUE TRUE TRUE
```

```
x>z ##Test if var > mad
```

```
##      C1  NO3  NH4  oP04  P04  Chla
```


```
## [1,] TRUE TRUE TRUE TRUE TRUE TRUE
```

What stands out most is that some chemicals have a rather large variance. Additionally each chemical seems to vary in mean quite significantly as well. 1c. The constant in the mad() function ensures consistency when calculating the median of a data set.

2abc. Altered histogram and boxplot of algae data set.


```
library(ggplot2)
ggplot(algae,aes(Chla,..density..))+
  geom_histogram(na.rm=T)+
  geom_density(na.rm = T)+
  geom_rug(aes(Chla),na.rm=T,inherit.aes = F)
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```




```
ggplot(algae,aes(season,a1))+
  geom_boxplot(na.rm=T,fill=c("brown2","green3","yellow","lightblue"))+
  labs(title="A conditioned Boxplot of Algae a1")
```

...A conditioned Boxplot of Algae a1...




2d.

```
ggplot(algae,aes(season,mn02))+  
  geom_boxplot(na.rm=T)
```



```
ggplot(algae,aes(season,C1))+  
  geom_boxplot(na.rm=T,fill=c("brown2","green3","yellow","lightblue"))
```



From the above boxplots we observe 2 seasonal outliers during the summer and winter when measuring mn02, and an additional 7 seasonal outliers when measuring Cl. Boxplots are a good way of finding outliers as they isolate the values that are significantly above or below the data set mean.

2e. Comparing mean v median and var v MAD

```
w1=algae%>%summarise_at(c(9,10),mean,na.rm=T)
w1##Mean of oP04 and P04
```

```
##      oP04      P04
## 1 75.24621 133.2588
```

```
x1=algae%>%summarise_at(c(9,10),var,na.rm=T)
x1##Var of oP04 and P04
```

```
##      oP04      P04
## 1 9421.302 18136.13
```

```
y1=algae%>%summarise_at(c(9,10),median,na.rm=T)
y1##Median of oP04 and P04
```

```
##      oP04      P04
## 1 35.928 85.95
```

```
z1=algae%>%summarise_at(c(9,10),mad,na.rm=T)
z1##MAD of oP04 and P04
```

```
##      oP04      P04
## 1 44.37125 102.9206
```

We clearly see that the mean and median, along with the variance and mad, differ greatly. This implies that

outliers are providing a rather significant skew when calculating the mean and variance of the two chemicals.

3a.Number of observations with missing values

```
na.test=is.na(algae)*1
colSums(na.test)
```

```
## season    size  speed  mxPH  mnO2    C1    NO3    NH4    oP04    P04
##      0      0      0      1      1      7      0      0      0      0
##   Chla    a1
##      8      0
```

We observe that 4 variables have missing values. mxPh and mnO2 have 1 missing value each, C1 has 7 missing values, and Chla has 8 missing values.

3b.Removing observations with missing values with filter()

```
algae.del=filter(algae,is.na(mxPH)==0&is.na(mnO2)==0&is.na(C1)==0&is.na(Chla)==0)
algae.cc=complete.cases(algae.del)
sum(algae.cc)
```

```
## [1] 169
```

Algae.del is a data set with 169 complete cases. 3c.Imputing unknowns with measures of central tendency

```
algae.mean=algae
for(j in c(1:180)){
  for(i in c(4:12)){
    pre.am=as.data.frame(algae[j,i])
    pre.algae.mean=mutate_if(pre.am,is.double,funs(ifelse(is.na(pre.am),mean(algae[,i],na.rm=T),algae[j,i])))
    algae.mean[j,i]=pre.algae.mean
  }
}
summarise(algae.mean, count=n())
```

```
## count
## 1 180
```

```
algae.mean[c(70,117,180),4:12]
```

```
##      mxPH mnO2    C1  NO3 NH4 oP04 P04    Chla  a1
## 70  6.6 11.3 40.2932 4.17 10    1    6 10.62003 47.1
## 117 6.6 10.8 40.2932 2.64 10    2   11 10.62003 46.9
## 180 5.7 10.8 40.2932 2.55 10    1    4 10.62003 16.8
```

3d. Imputing unknowns using correlation

```
cor(algae[,4:12],use="complete.obs")
```

```
##      mxPH    mnO2    C1    NO3    NH4
## mxPH  1.00000000 -0.02913045  0.1557075  0.03339517 -0.12005457
## mnO2 -0.02913045  1.00000000 -0.3348301 -0.03292011 -0.28278234
## C1    0.15570753 -0.33483005  1.0000000  0.45590366  0.16225274
## NO3   0.03339517 -0.03292011  0.4559037  1.00000000  0.14587533
## NH4   -0.12005457 -0.28278234  0.1622527  0.14587533  1.00000000
## oP04  0.04653433 -0.39999298  0.3968281  0.30279529  0.56762718
## P04   0.04409164 -0.48414176  0.4710502  0.32852261  0.62113218
## Chla  0.43879072 -0.16058108  0.1677026  0.06764031 -0.03901627
## a1    -0.23274670  0.27512691 -0.4145950 -0.37398729 -0.17565791
##      oP04    P04    Chla    a1
```

```
## mxPH 0.04653433 0.04409164 0.43879072 -0.2327467
## mnO2 -0.39999298 -0.48414176 -0.16058108 0.2751269
## Cl 0.39682805 0.47105019 0.16770264 -0.4145950
## NO3 0.30279529 0.32852261 0.06764031 -0.3739873
## NH4 0.56762718 0.62113218 -0.03901627 -0.1756579
## oPO4 1.00000000 0.93253518 0.04982285 -0.4200985
## PO4 0.93253518 1.00000000 0.15925064 -0.4508198
## Chla 0.04982285 0.15925064 1.00000000 -0.2925959
## a1 -0.42009846 -0.45081979 -0.29259589 1.0000000
```

```
fit1=lm(Chla~mxPH,data=algae)
coefficients(fit1)
```

```
## (Intercept)          mxPH
## -118.88773      16.04646
```

```
predict(fit1,data.frame(mxPH=c(5.7,6.5,6.6,7.83,9.7)))
```

```
##          1          2          3          4          5
## -27.422907 -14.585738 -12.981092   6.756055  36.762936
```

#Obtained values when mxPH is measured to be 5.7,6.5,6.6,7.83,and 9.7

These are odd numbers to obtain since some of them are negative.

4a. Partitioning of data set algae.mean

```
set.seed(10)
cut.id=cut(1:180,6,labels = F)
cut.id.rand=sample(cut.id,180)
algae.sample.id=cut(cut.id.rand,6,labels=F)
```

4b. 6 fold cross validation

```
library(plyr)
```

```
## -----
## You have loaded plyr after dplyr - this is likely to cause problems.
## If you need functions from both plyr and dplyr, please load plyr first, then dplyr:
## library(plyr); library(dplyr)
```

```
## -----
```

```
##
```

```
## Attaching package: 'plyr'
```

```
## The following objects are masked from 'package:dplyr':
```

```
##
```

```
##      arrange, count, desc, failwith, id, mutate, rename, summarise,
##      summarize
```

```
do.chunk <- function(chunkid, chunkdef, dat){ # function argument
train = (chunkdef != chunkid)
Ytr = dat[train,]$a1 # get true response values in training set
Yvl = dat[!train,]$a1 # get true response values in validation set
lm.a1 <- lm(a1~., data = dat[train,])
predYtr = predict(lm.a1) # predict training response values
predYvl = predict(lm.a1,dat[!train,]) # predict validation values
data.frame(fold = chunkid,
train.error = mean((predYtr - Ytr)^2), # compute and store training errors
```

```
val.error = mean((predYv1 - Yv1)^2)) # compute and store validation errors
}
cv.6f=ldply(1:6,do.chunk,algae.sample.id,algae.mean)
cv.6f
```

```
##   fold train.error val.error
## 1    1    231.8446  466.6458
## 2    2    267.3170  264.8416
## 3    3    260.2019  295.8210
## 4    4    247.8380  383.4553
## 5    5    270.7286  247.7628
## 6    6    274.1990  233.6241
```

5. Testing algae.test file.

```
algaeTest = read.table("algae-test.txt", header=T, na.strings="NA")
```

```
val.error.avg=mean(cv.6f[,3])
val.error.avg##average of valuation error
```

```
## [1] 315.3584
```

```
lm.a1.2=lm(a1~.,data=algae.mean)
```

```
a1.test.pred=predict(lm.a1.2,algaeTest)
a1.test.pred
```

```
##          1          2          3          4          5          6
##  7.4336194  0.8124484 15.1236371 21.7018260 25.4605555 35.8558252
##          7          8          9         10         11         12
## 28.5927301 34.1002476 34.4933167 29.1322886          NA 26.0753508
##         13         14         15         16         17         18
## 41.5008572 23.8306777 37.9088291 14.4221764 10.1389309 26.7856776
##         19         20         21         22         23         24
## 28.0539124 31.3413817 28.3401972 30.3474327          NA          NA
##         25         26         27         28         29         30
## 36.9906760 40.0117574 38.0789920 15.2655612  8.5607639 24.9025220
##         31         32         33         34         35         36
## 21.9392482 25.2630485 24.1665620 19.3906134 15.8554695 18.5614626
##         37         38         39         40         41         42
## -6.3539289 -4.3988611  8.6456563 -6.2158181  6.0432748 17.2266901
##         43         44         45         46         47         48
## -1.5874397 18.5975122 -6.5439746 15.9606574 14.1833677 23.3995388
##         49         50         51         52         53         54
## 23.5475449 29.4925914  3.0469893 26.4815484 20.7319677 -7.7054715
##         55         56         57         58         59         60
## -12.6053118 23.6375182 22.5877617 19.2731724 20.2494723 10.1283481
##         61         62         63         64         65         66
## 16.8493796 -34.3547492 16.6198052 -4.8191644  2.1351614 -10.0340883
##         67         68         69         70         71         72
## -8.4101419 22.1874309 15.3240174  4.5598152 11.4773656 -3.3972793
##         73         74         75         76         77         78
## 25.7980437 25.9591162 25.3244379          NA 25.1336833 -1.1539189
##         79         80         81         82         83         84
## 15.2409406 15.6371562 12.8505710 19.5443274 12.5991842 16.7327584
##         85         86         87         88         89         90
```



```
## 3.6777198      NA 7.7334436 17.6224599 23.1728851 36.7552461
## 91      92      93      94      95      96
## -54.6861243 38.4101709 37.9127800 29.8926220 35.3070955 10.7771792
## 97      98      99     100     101     102
## 14.2536802 -1.3950356 -6.0126564 24.9235553 43.4589283 32.7516646
## 103     104     105     106     107     108
## 35.3140973 40.3927125 39.2078702 36.7025597 38.7128532 35.9748090
## 109     110     111     112     113     114
## 39.0257476 19.0012175 19.3287288 8.4565746 7.7849182 13.4386025
## 115     116     117     118     119     120
## 0.9875589 4.7262051 16.1004242 1.6524832 6.2715781 6.1175048
## 121     122     123     124     125     126
## 17.9821767 17.7405451 25.9169313 24.0959127 23.7219232 28.7406761
## 127     128     129     130     131     132
## 5.0717693 2.3554747 1.9753651 21.5741983 22.2987850 20.7516918
## 133     134     135     136     137     138
## 4.3060994 6.9549463 -1.9994972 12.7915820 15.3610764 18.4532987
## 139     140     141     142     143     144
## 17.1504364 14.8373311 10.3433198 14.5213188 15.8886815 -2.0861936
## 145     146     147     148     149     150
## 0.6774467 -7.7484079 0.3642028 25.5103241 14.3845324 18.0971744
## 151     152     153     154     155     156
## 14.5800206 13.6600917 30.1768521 35.1738197 35.1270354 23.1504807
## 157     158     159     160
## 38.5180005 25.2949296 25.0570888 16.2703687
```

```
train.error.test=mean((a1.test.pred-algaeTest$a1)^2,na.rm=T)
train.error.test##test error from algaeTest
```

```
## [1] 290.6248
```

```
val.error.avg##Test error from 6 fold C.V
```

```
## [1] 315.3584
```


The test error of the two data sets are rather similar but still vary slightly. Considering the number of predictor variables, their possible interactions, and total observations, this result is expected.

6a.

```
library(ISLR)
data(Carseats)
attach(Carseats)
```

```
ggplot(Carseats,aes(Income,Sales))+
  geom_point()+
  geom_smooth()
```

```
## 'geom_smooth()' using method = 'loess'
```



From this plot there does not appear to be a significantly large relation between sales and income. This matches my intuition as I believe child car seats are long lasting and are often reused between children. I believe carseats are seldomly repeatedly purchased by individual customers unless required. I assume carseat sales have more to do with measurements related to the number of children in a given area than it does with income.

6b. Fitting linear models to the p-th degree of Income and running a 6 fold C.V

```
fit.sales.inc=lm(Sales~poly(Income,10,row=F),data=Carseats)
summary(fit.sales.inc)
```

```
##
## Call:
## lm(formula = Sales ~ poly(Income, 10, raw = F), data = Carseats)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-7.0727	-1.9203	-0.1168	1.7017	8.9917

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	7.49633	0.14026	53.445	<2e-16 ***
poly(Income, 10, raw = F)1	8.57181	2.80527	3.056	0.0024 **
poly(Income, 10, raw = F)2	-1.90088	2.80527	-0.678	0.4984
poly(Income, 10, raw = F)3	-0.22592	2.80527	-0.081	0.9359
poly(Income, 10, raw = F)4	-0.45907	2.80527	-0.164	0.8701
poly(Income, 10, raw = F)5	1.39137	2.80527	0.496	0.6202
poly(Income, 10, raw = F)6	-3.89104	2.80527	-1.387	0.1662

```
## poly(Income, 10, raw = F)7    2.04818    2.80527    0.730    0.4658
## poly(Income, 10, raw = F)8    4.39531    2.80527    1.567    0.1180
## poly(Income, 10, raw = F)9    1.75371    2.80527    0.625    0.5322
## poly(Income, 10, raw = F)10 -0.07317    2.80527   -0.026    0.9792
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.805 on 389 degrees of freedom
## Multiple R-squared:  0.03803,    Adjusted R-squared:  0.0133
## F-statistic: 1.538 on 10 and 389 DF,  p-value: 0.1237

set.seed(10)
cut.id.5=cut(1:400,6,labels = F)
cut.id.rand.5=sample(cut.id.5,400)
carseats.sample.id=cut(cut.id.rand.5,6,labels=F)
carseats.mut=mutate(Carseats,SampleID=carseats.sample.id)
```

6 Fold Cross Validation Function

```
do.chunk5 <- function(chunkid, chunkdef, dat, p){# function argument
train = (chunkdef != chunkid)
res = data.frame(degree=integer(), fold=integer(),
train.error=double(), val.error=double())
if (p==0) {
## Your code here
Ytr.5=dat[train,]$Sales
Yvl.5=dat[!train,]$Sales
## Fit an intercept only model to the data.
## Using poly(Income, degree=0) will cause an error
sales.incp=lm(Sales~1,data=dat[train,])
PYtr.5=predict(sales.incp,dat[train,])
PYvl.5=predict(sales.incp,dat[!train,])
## Update residual
res = data.frame(degree=p, fold=chunkid,
train.error=mean((PYtr.5-Ytr.5)^2),
val.error=mean((PYvl.5-Yvl.5)^2))
res
} else {
## Your code here
Ytr.5=dat[train,]$Sales
Yvl.5=dat[!train,]$Sales
## Fit a polynomial regression or order p.
## Use poly(Income, degree=p, raw=FALSE)
sales.incp=lm(Sales~poly(Income,p),data=dat[train,])
PYtr.5=predict(sales.incp,dat[train,])
PYvl.5=predict(sales.incp,dat[!train,])
## Update residual
res = data.frame(degree=p, fold=chunkid,
train.error=mean((PYtr.5-Ytr.5)^2),
val.error=mean((PYvl.5-Yvl.5)^2))
res
}
}
```

Returning each CV of carseats data set

```

cvlist=c()
mve=rep(0,11)
mte=rep(0,11)
for(i in c(0:10)){
  cv.6f.carseats=ldply(1:6,do.chunk5,carseats.sample.id,Carseats,i)
  mve[i+1]=mean(cv.6f.carseats$val.error)
  mte[i+1]=mean(cv.6f.carseats$train.error)
  print(cv.6f.carseats)
  cvlist=append(cvlist,cv.6f.carseats)
}

```

```

## degree fold train.error val.error
## 1 0 1 8.518380 5.163070
## 2 0 2 7.966481 7.902722
## 3 0 3 7.756241 8.989137
## 4 0 4 7.526500 10.095775
## 5 0 5 7.809419 8.699561
## 6 0 6 8.143375 7.071808
## degree fold train.error val.error
## 1 1 1 8.373222 4.798782
## 2 1 2 7.782577 7.719448
## 3 1 3 7.524387 9.073789
## 4 1 4 7.317815 10.035983
## 5 1 5 7.700618 8.195184
## 6 1 6 7.901821 7.189609
## degree fold train.error val.error
## 1 2 1 8.372869 4.781148
## 2 2 2 7.762206 7.781148
## 3 2 3 7.495361 9.196735
## 4 2 4 7.290461 10.148359
## 5 2 5 7.695717 8.167896
## 6 2 6 7.901820 7.190583
## degree fold train.error val.error
## 1 3 1 8.371449 4.800467
## 2 3 2 7.759393 7.803331
## 3 3 3 7.495351 9.196521
## 4 3 4 7.290161 10.153321
## 5 3 5 7.687124 8.255757
## 6 3 6 7.901481 7.197431
## degree fold train.error val.error
## 1 4 1 8.368867 4.816662
## 2 4 2 7.757257 7.813426
## 3 4 3 7.495269 9.195032
## 4 4 4 7.278740 10.307075
## 5 4 5 7.683365 8.278994
## 6 4 6 7.896173 7.230387
## degree fold train.error val.error
## 1 5 1 8.366804 4.803476
## 2 5 2 7.744634 7.861714
## 3 5 3 7.493632 9.177620
## 4 5 4 7.266668 10.360326
## 5 5 5 7.665263 8.358911
## 6 5 6 7.895677 7.255280
## degree fold train.error val.error

```

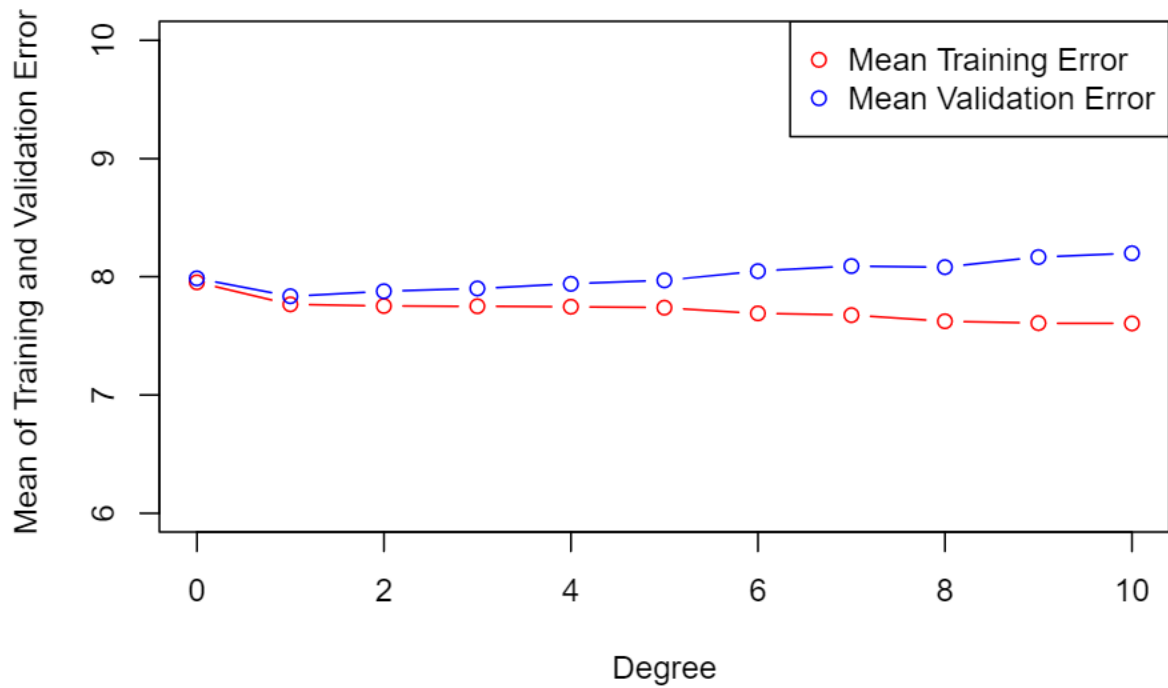
```
## 1      6      1      8.314426  4.851425
## 2      6      2      7.668704  8.055806
## 3      6      3      7.462140  9.111067
## 4      6      4      7.266611 10.379073
## 5      6      5      7.590881  8.568181
## 6      6      6      7.842271  7.318575
## degree fold train.error val.error
## 1      7      1      8.297903  4.881005
## 2      7      2      7.622620  8.302501
## 3      7      3      7.455335  9.086806
## 4      7      4      7.250783 10.402811
## 5      7      5      7.590860  8.574298
## 6      7      6      7.835940  7.296061
## degree fold train.error val.error
## 1      8      1      8.249949  4.822586
## 2      8      2      7.502641  8.716494
## 3      8      3      7.426507  8.950169
## 4      8      4      7.227724 10.261017
## 5      8      5      7.547281  8.486072
## 6      8      6      7.788168  7.251606
## degree fold train.error val.error
## 1      9      1      8.240748  4.823584
## 2      9      2      7.492778  8.715575
## 3      9      3      7.425849  8.930280
## 4      9      4      7.191047 10.453235
## 5      9      5      7.539225  8.705025
## 6      9      6      7.759523  7.374172
## degree fold train.error val.error
## 1     10      1      8.240486  4.825421
## 2     10      2      7.492546  8.719564
## 3     10      3      7.423976  8.953799
## 4     10      4      7.190618 10.457062
## 5     10      5      7.528974  8.822139
## 6     10      6      7.755880  7.418140
```

Couldnt figure out how to make the ldply function calculate different values of p in addition to chunkid. Eventually I nested it in a for loop to get the required results.

Plotting average Traing and Validation error

```
plot(c(0:10),mte,col="red",type="b",ylim = c(6,10),
     main="Plot of Training and Validation Error",xlab="Degree",
     ylab=" Mean of Training and Validation Error")
lines(c(0:10),mve,col="blue",type="b")
legend("topright",legend = c("Mean Training Error","Mean Validation Error"),pch=c(1,1),col=c("red"
```

Plot of Training and Validation Error



From this plot I see that the training error related to this data set decreases, and validation error increases, slightly with higher degree polynomials.

Based on the graph, I would choose the first degree of the polynomial. The reason being that it has the smallest average validation error.